

# Deep Dive on AWS Lambda

Capabilities, benefits, and best practices

## Presenter

Vyom Nagrani

Manager Product Management, AWS Lambda

## Q&A Moderators

Stefano Buliani, Solution Architect, AWS Serverless

Vishal Singh, Sr. Product Manager, AWS Lambda

Orr Weinstein, Sr. Product Manager, AWS Lambda

January 18<sup>th</sup>, 2017

# Agenda for today's Webinar

- Working with AWS Lambda
- Development and testing on AWS Lambda
- Deployment and ALM for AWS Lambda
- Security and scaling on AWS Lambda
- Debugging and operations for AWS Lambda
- Questions & answers

# Working with AWS Lambda

# Working with AWS Lambda

## EVENT SOURCE



Changes in  
data state



Requests to  
endpoints



Changes in  
resource state



## FUNCTION



Node  
Python  
Java  
C#

## SERVICES (ANYTHING)



# Benefits of AWS Lambda

Productivity-focused compute platform to build powerful, dynamic, modular applications in the cloud

**1**

**No infrastructure  
to manage**



Focus on business logic

**2**

**Cost-effective and  
efficient**



Pay only for what you use

**3**

**Bring your  
own code**



Run code in standard  
languages

# Event sources that trigger AWS Lambda

## DATA STORES



Amazon  
S3



Amazon  
DynamoDB



Amazon  
Kinesis



Amazon  
Cognito

## ENDPOINTS



Amazon  
Alexa



Amazon  
API Gateway



AWS  
IoT

## CONFIGURATION REPOSITORIES



AWS  
CloudFormation



AWS  
CloudTrail



AWS  
CodeCommit



Amazon  
CloudWatch

## EVENT/MESSAGE SERVICES



Amazon  
SES



Amazon  
SNS



Cron events

*... and the list will continue to grow!*

# Key scenarios and use cases for AWS Lambda



## Data processing

Stateless processing of discrete or streaming updates to your data-store or message bus



## App backend development

Execute server side backend logic for web, mobile, device, or voice user interactions



## Control systems

Customize responses and response workflows to state and data changes within AWS

# Development and testing on AWS Lambda



# Getting started on AWS Lambda



## Bring your own code

- Node.js 4.3, Java 8, Python 2.7, C#




## Simple resource model

- Select power rating from 128 MB to 1.5 GB
- CPU and network allocated proportionately



## Flexible use

- Synchronous or asynchronous
- Integrated with other AWS services



## Stateless

- Persist data using external storage
- No affinity or access to underlying infrastructure

# Anatomy of a Lambda function

## Handler() function

- The method in your code where AWS Lambda begins execution

## Event object

- Pre-defined object format for AWS integrations & events
- Java & C# support simple data types, POJOs/POCOs, and Stream input/output

## Context object

- Use methods and properties like `getRemainingTimeInMillis()`, `identity`, `awsRequestId`, `invokedFunctionArn`, `clientContext`, `logStreamName`

# FunctionConfiguration metadata

## VpcConfig

- Enables private communication with other resources within your VPC
- Provide EC2 security group and subnets, auto-creates ENIs
- Internet access can be added though NAT Gateway

## DeadLetterConfig

- Failed events sent to your SQS queue / SNS topic
- Redrive messages that Lambda could not process
- Currently available for asynchronous invocations only

## Environment

- Add custom key/value pairs as part of configuration
- Reuse code across different setups or passwords
- Encrypted with specified KMS key on server, decrypted at container init

# AWS Lambda limits

Resource Limits	Default Limit
Ephemeral disk capacity ("/tmp" space)	512 MB
Number of file descriptors	1024
Number of processes and threads (combined total)	1024
Maximum execution duration per request	300 seconds
Invoke request body payload size (RequestResponse)	6 MB
Invoke request body payload size (Event)	128 K
Invoke response body payload size (RequestResponse)	6 MB
Dead-letter payload size (Event)	128 K
Deployment Limits	Default Limit
Lambda function deployment package size (.zip/.jar file)	50 MB
Size of code/dependencies that you can zip into a deployment package (uncompressed zip/jar size)	250 MB
Total size of all the deployment packages that can be uploaded per region	75 GB
Total size of environment variables set	4 KB
Throttling Limits (can request service limit increase)	Default Limit
Concurrent executions	100

# The container model

## Container reuse

- Declarations in your Lambda function code outside handler()
- Disk content in /tmp
- Background processes or callbacks
- Make use of container reuse opportunistically, e.g.
  - Load additional libraries
  - Cache static data
  - Database connections

## Cold starts

- Time to set up a new container and do necessary bootstrapping when a Lambda function is invoked for the first time or after it has been updated
- Ways to reduce cold start latency
  - More memory = faster performance, lower start up time
  - Smaller function ZIP loads faster
  - Node.js and Python start execution faster than Java and C#

# The execution environment

## Underlying OS

- Public Amazon Linux AMI version (*amzn-ami-hvm-2016.03.3.x86\_64-gp2*)
- Linux kernel version (*4.4.23-31.54.amzn1.x86\_64*)
- Compile native binaries against this environment – can be used to bring your own runtime!
- Changes over time, always check the latest versions supported [here](#)

## Available libraries

- ImageMagick (nodejs wrapper and native binary)
- OpenJDK 1.8, .NET Core 1.0.1
- AWS SDK for JavaScript version 2.6.9
- AWS SDK for Python (Boto 3) version 1.4.1, Botocore version 1.4.61
- Embed your own SDK/libraries if you depend on a specific version

# Deployment and ALM for AWS Lambda

# Building a deployment package

## Node.js & Python

- .zip file consisting of your code and any dependencies
- Use npm/pip to install libraries
- All dependencies must be at root level

## Java

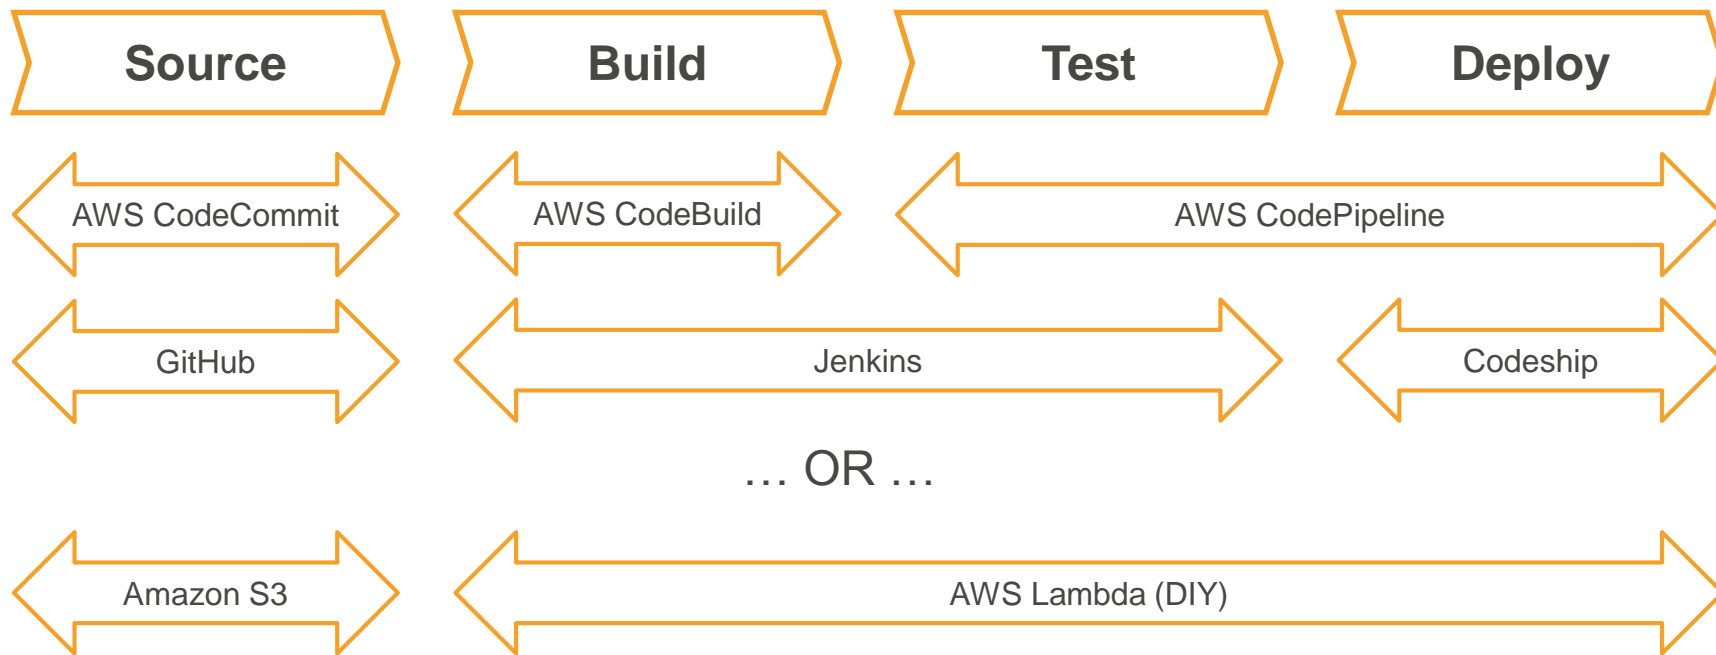
- Either .zip file with all code/dependencies, or standalone .jar
- Use Maven / Eclipse IDE plugins
- Compiled class & resource files at root level, required jars in /lib directory

## C# (.NET Core)

- Either .zip file with all code/dependencies, or a standalone .dll
- Use Nuget / VisualStudio plugins
- All assemblies (.dll) at root level



# Managing continuous delivery



# Deployment tools and frameworks available

## CloudFormation

- AWS Serverless Application Model - extension optimized for Serverless
- New Serverless resources – APIs, Functions, Tables
- Open specification (Apache 2.0)

## Chalice

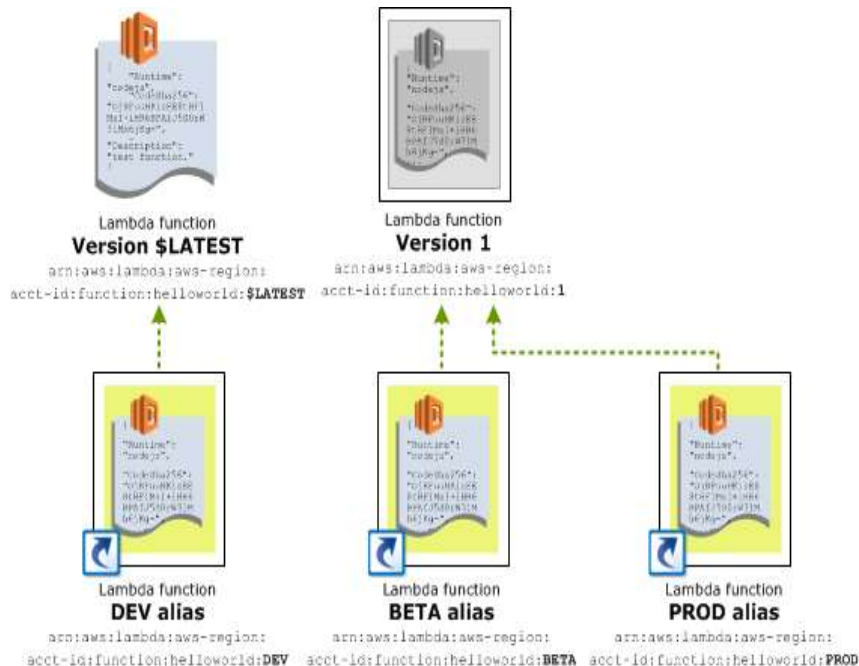
- Python serverless micro-framework
- Quickly create and deploy applications
- Set up AWS Lambda and Amazon API Gateway endpoint
- <https://github.com/awslabs/chalice>

## Third-party tools

- Serverless Framework (<https://serverless.com/>)
- Apex Serverless Architecture (<http://apex.run/>)
- DEEP Framework by Mitoc Group (<https://github.com/MitocGroup/deep-framework>)

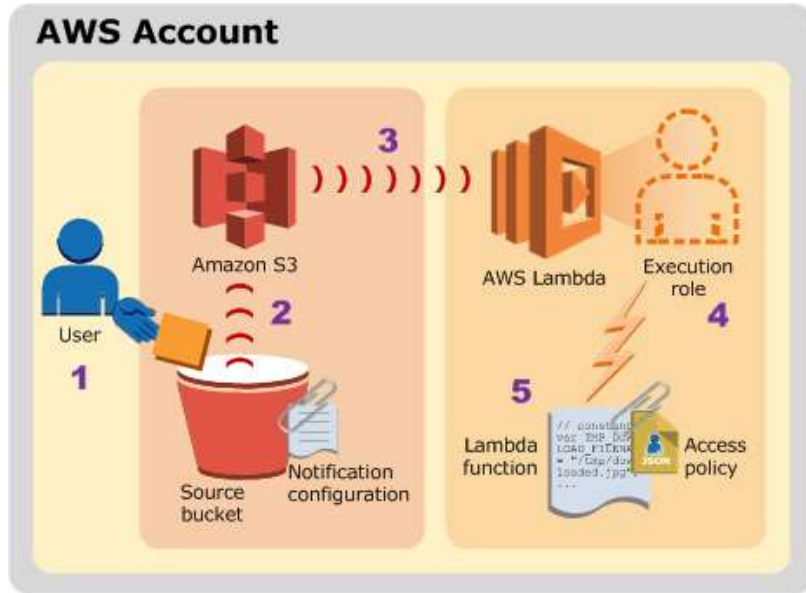
# Function versioning and aliases

- Versions = immutable copies of code + configuration
- Aliases = mutable pointers to versions
- Development against \$LATEST version
- Each version/alias gets its own ARN
- Enables rollbacks, staged promotions, “locked” behavior for client



# Security and scaling on AWS Lambda

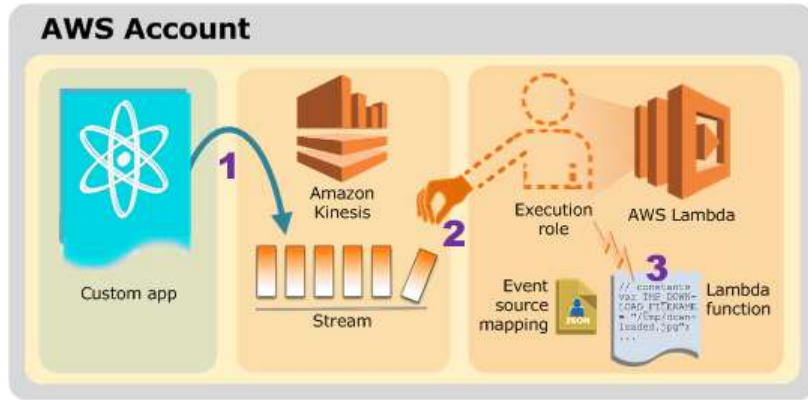
# The push model and resource policies



## Function (resource) policy

- Permissions you grant to your Lambda function determine which service or event source can invoke your function
- Resource policies make it easy to grant cross-account permissions to invoke your Lambda function

# The pull model and IAM roles



## IAM (execution) role

- Permissions you grant to this role determine what your AWS Lambda function can do
- If event source is Amazon DynamoDB or Amazon Kinesis, then add read permissions in IAM role

# Concurrent executions and throttling

## Determining concurrency

- For stream-based event sources:  
Number of shards per stream is the unit of concurrency
- For all other event sources: Request rate and duration drives concurrency ( $concurrency = requests\ per\ second * function\ duration$ )

## Throttle behavior

- For stream-based event sources:  
Automatically retried until data expires
- For Asynchronous invocations:  
Automatically retried for up to six hours, with delays between retries
- For Synchronous invocations: Invoking application receives a 429 error and is responsible for retries

# Other scaling considerations

## For Lambda

- Remember, a throttle is NOT an error!
- If you expect sudden large spikes in demand, consider Asynchronous invocations to Lambda
- Proactively engage AWS Support to increase your throttling limits

## For upstream/downstream services

- Build retries/backoff in client applications and upstream setup
- Make sure your downstream setup “keeps up” with Lambda scaling
- Limit concurrency when connecting to relational databases



# Debugging and operations for AWS Lambda

# Errors and retries

## Types of errors

- 4xx Client Error: Can be fixed by developer, e.g. InvalidParameterValue (400), ResourceNotFound (404), RequestTooLarge (413), etc.
- 5xx Server Error: Most can be fixed by admin, e.g. EC2 ENI management errors (502)

## Retry policy

- For stream-based event sources: Automatically retried until data expires
- For Asynchronous invocations: Automatically retried 2 extra times, then published to dead-letter queue
- For Synchronous invocations: Invoking application receives an error code and is responsible for retries

# Tracing and tracking



COMING  
SOON!

## Integration with AWS X-Ray

- Collects data about requests that your application serves
- Visibility into the AWS Lambda service (dwell time, number of retries, latency and errors)
- Detailed breakdown of your function's performance, including calls made to downstream services and endpoints

## Integration with AWS CloudTrail

- Captures calls made to AWS Lambda API; delivers log files to Amazon S3
- Tracks the request made to AWS Lambda, the source IP address from which the request was made, who made the request, when it was made
- All control plane APIs can be tracked (no versioning/aliasing and invoke API)

# Troubleshooting and monitoring

## Logs

- Every invocation generates START, END, and REPORT entries in CloudWatch Logs
- User logs included
  - Node.js – `console.log()`, `console.error()`, `console.warn()`, `console.info()`
  - Java – `log4j.*`, `LambdaLogger.log()`, `system.out()`, `system.err()`
  - Python – `print`, `logging.*`
  - C# – `LambdaLogger.Log()`, `ILambdaContext.Logger.Log()`, `console.write()`, `console.writeline()`

## Metrics

- Default (Free) Metrics: Invocations, Duration, Throttles, Errors – available as CloudWatch Metrics
- Additional Metrics: Create custom metrics for tracking health/status
  - Function code vs log-filters
  - Ops-centric vs. business-centric

# Conclusion and next steps

## Key takeaway

AWS Lambda is one of the core components of the platform AWS provides to develop serverless applications

## Next steps

1. Stay up to date with AWS Lambda on the [Compute blog](#) and check out our [detail page](#) for more scenarios.
2. Send us your questions, comments, and feedback on the [AWS Lambda Forums](#).

# Thank you!